

Electronic and Telecommunication engineering
University of Moratuwa



EN3160 - Image Processing and Machine Vision
A01 Intensity Transformations and Neighborhood Filtering

RANAVIRAJA R.W.H.M.T.A. - Index No. 200520X

GitHub: [GitHub](#)

20th August, 2023

Q1.


```
c = np.array([(100, 50), (150, 255)])

t1 = np.linspace(0, c[0, 1], c[0, 1] + 1 - 0).astype('uint8')
print(len(t1))

t2 = np.linspace(c[0, 0], c[1, 1], c[1,0] - c[0, 1]).astype('uint8')
print(len(t2))

t3 = np.linspace(c[1, 0], c[1, 1], c[1, 1] - c[1, 0]).astype('uint8')
print(len(t3))

transform = np.concatenate((t1, t2, t3), axis=0).astype('uint8')
print(len(transform))
```



Original Image



Transformed Image




Q2.


(a) White matter

```
t1 = np.linspace(0, 0, 181).astype('uint8')
t2 = np.linspace(200, 255, 255-180).astype('uint8')
transform_whitematter = np.concatenate((t1, t2), axis=0).astype('uint8')
```

```
# plot transformed white matter image
image_whitematter = cv.LUT(img_orig, transform_whitematter)
```



Original Image



White Matter Image


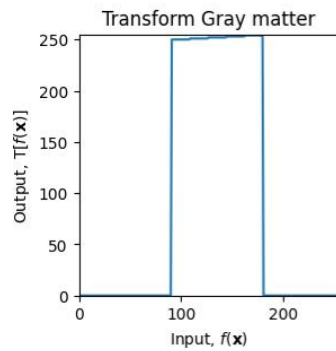


(b) Gray matter




```
t1 = np.linspace(0,0, 91).astype('uint8')
t2 = np.linspace(250, 255, 180-90).astype('uint8')
t3 = np.linspace(0, 0, 255-180).astype('uint8')
transform_graymatter = np.concatenate((t1, t2,t3), axis=0).astype('uint8')
```

```
# plot transformed white matter image
```

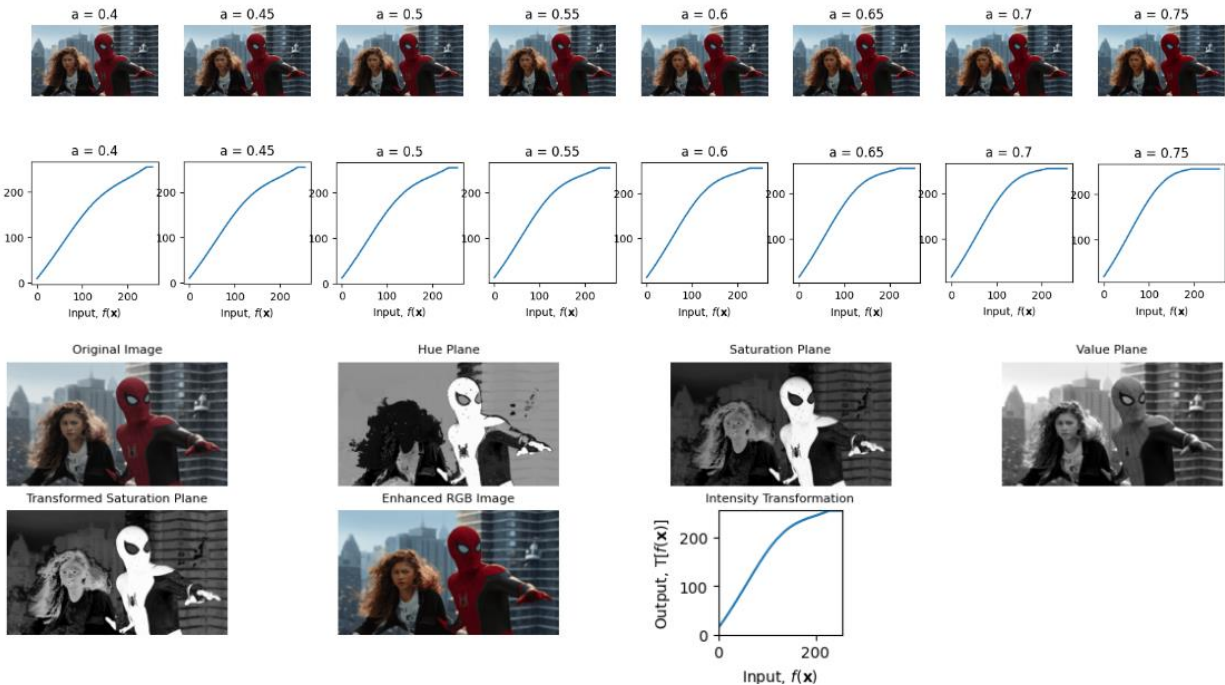
```
image_graymatter = cv.LUT(img_orig, transform_graymatter)
```



Q3. Selected gamma value = 0.65



Q4. Selected 'a' value =0.6



Q5.

```
# Calculate histogram of the original image
histogram = np.zeros((256,), dtype=np.uint16)
for i in range(rows):
    for j in range(cols):
        intensity = input_image[i, j]
        histogram[intensity] += 1

# Calculate cumulative distribution function (CDF)
cdf = np.zeros((256,), dtype=np.uint16)
for i in range(256):
    for j in range(i + 1):
        cdf[i] += histogram[j] * (255 / (rows * cols))
    cdf[i] = round(cdf[i], 0)
cdf = cdf.astype(np.uint16)

# Apply histogram equalization
output_image = np.zeros_like(input_image)
for i in range(rows):
    for j in range(cols):
        intensity = input_image[i, j]
        output_image[i, j] = cdf[intensity]
```



Q6.

```
# Thresholding the saturation plane to create a foreground mask
_, foreground_mask = cv.threshold(sat, 15, 255, cv.THRESH_BINARY)
# foreground_mask = cv.morphologyEx(foreground_mask, cv.MORPH_CLOSE, cv.getStructuringElement(cv.MORPH_ELLIPSE, (80,80) ))

plt.subplot(2, 5, 5)
plt.imshow(foreground_mask, cmap='gray')
plt.title("Foreground Mask")
plt.axis('off')

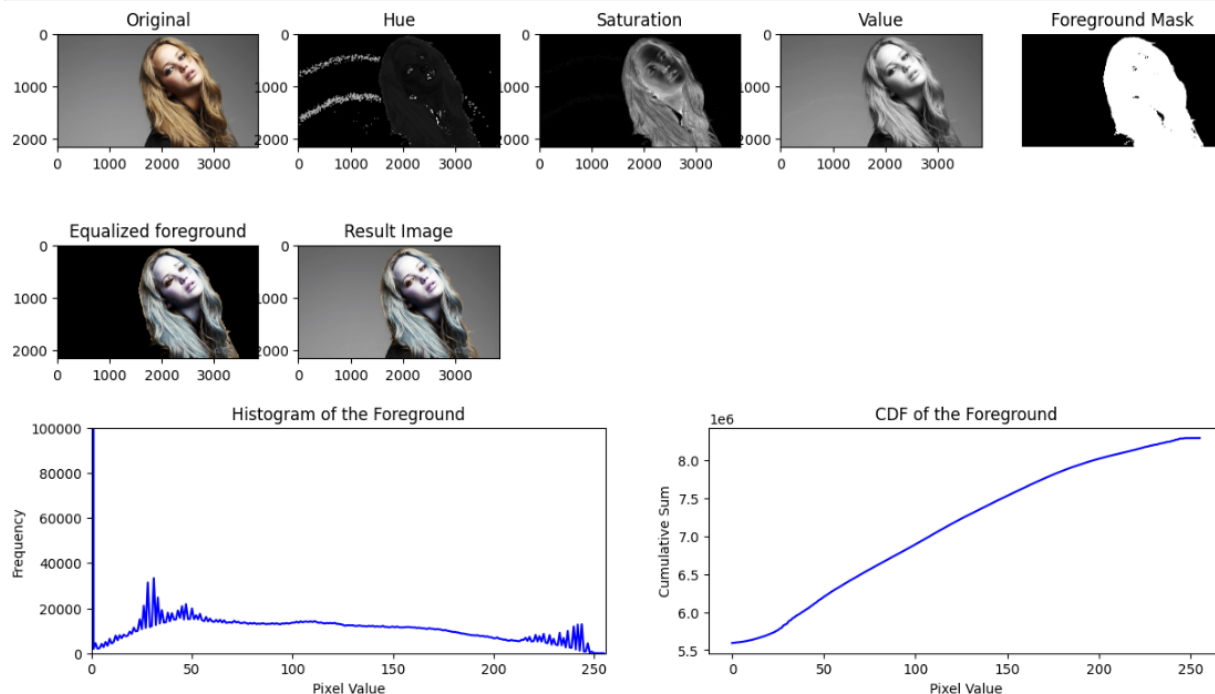
# Apply bitwise AND operation to get the foreground only
foreground = cv.bitwise_and(img, img, mask=foreground_mask)

# Convert foreground to grayscale
foreground_gray = cv.cvtColor(foreground, cv.COLOR_BGR2GRAY)

# Compute the histogram of the foreground
hist_foreground = cv.calcHist([foreground_gray], [0], None, [256], [0, 256])


# Histogram-equalize the foreground channels separately
equalized_r = cv.equalizeHist(foreground[:, :, 0])
equalized_g = cv.equalizeHist(foreground[:, :, 1])
equalized_b = cv.equalizeHist(foreground[:, :, 2])

equalized_foreground = cv.merge((equalized_r, equalized_g, equalized_b))
```



Q7.

```
# write a function to filter the given image with given kernel
def filterImg(image, kernel):
    img_h, img_w = image.shape
    kernel_h, kernel_w = kernel.shape
    pad_h = kernel_h // 2
    pad_w = kernel_w // 2
    image_float = cv.normalize(image.astype('float'), None, 0.0, 1.0, cv.NORM_MINMAX)
    output = np.zeros(image.shape, 'float')
    for y in range(pad_h, img_h-pad_h):
        for x in range(pad_w, img_w-pad_w):
            output[y, x] = np.dot(image_float[y - pad_h:y + pad_h + 1, x - pad_w:x + pad_w + 1].flatten(), kernel.flatten())
    return output
```



Q8.

```
def zoom_nearest_neighbor(image, scale_factor):
    h, w = image.shape[:2]
    new_h = int(h * scale_factor)
    new_w = int(w * scale_factor)
    zoomed_image = cv2.resize(image, (new_w, new_h), interpolation=cv2.INTER_NEAREST)
    return zoomed_image

def zoom_bilinear(image, scale_factor):
    h, w = image.shape[:2]
    new_h = int(h * scale_factor)
    new_w = int(w * scale_factor)
    zoomed_image = cv2.resize(image, (new_w, new_h), interpolation=cv2.INTER_LINEAR)
    return zoomed_image

def normalized_ssd(image1, image2):
    diff = image1.astype(np.float32) - image2.astype(np.float32)
    ssd = np.sum(diff**2)
    normalized_ssd = ssd / (image1.shape[0] * image1.shape[1])
    return normalized_ssd
```

Image 01:	Normalized SSD (Nearest Neighbor): 71.4944	Normalized SSD (Bilinear Interpolation): 71.6791
Image 02:	Normalized SSD (Nearest Neighbor): 10.2422	Normalized SSD (Bilinear Interpolation): 10.3817
Image 03:	Normalized SSD (Nearest Neighbor): 75.0214	Normalized SSD (Bilinear Interpolation): 57.5857
Image 04:	Normalized SSD (Nearest Neighbor): 593.1880	Normalized SSD (Bilinear Interpolation): 593.3768
Image 05:	Normalized SSD (Nearest Neighbor): 235.7569	Normalized SSD (Bilinear Interpolation): 235.9187
Image 06:	Normalized SSD (Nearest Neighbor): 118.1741	Normalized SSD (Bilinear Interpolation): 118.2930
Image 07:	Normalized SSD (Nearest Neighbor): 143.8783	Normalized SSD (Bilinear Interpolation): 144.0937
Image 08:	Normalized SSD (Nearest Neighbor): 26.2169	Normalized SSD (Bilinear Interpolation): 20.7591
Image 09:	Normalized SSD (Nearest Neighbor): 20.0270	Normalized SSD (Bilinear Interpolation): 20.1700
Image 10:	Normalized SSD (Nearest Neighbor): 135.7616	Normalized SSD (Bilinear Interpolation): 117.7240
Image 11:	Normalized SSD (Nearest Neighbor): 168.3602	Normalized SSD (Bilinear Interpolation): 92.9710

Q9.

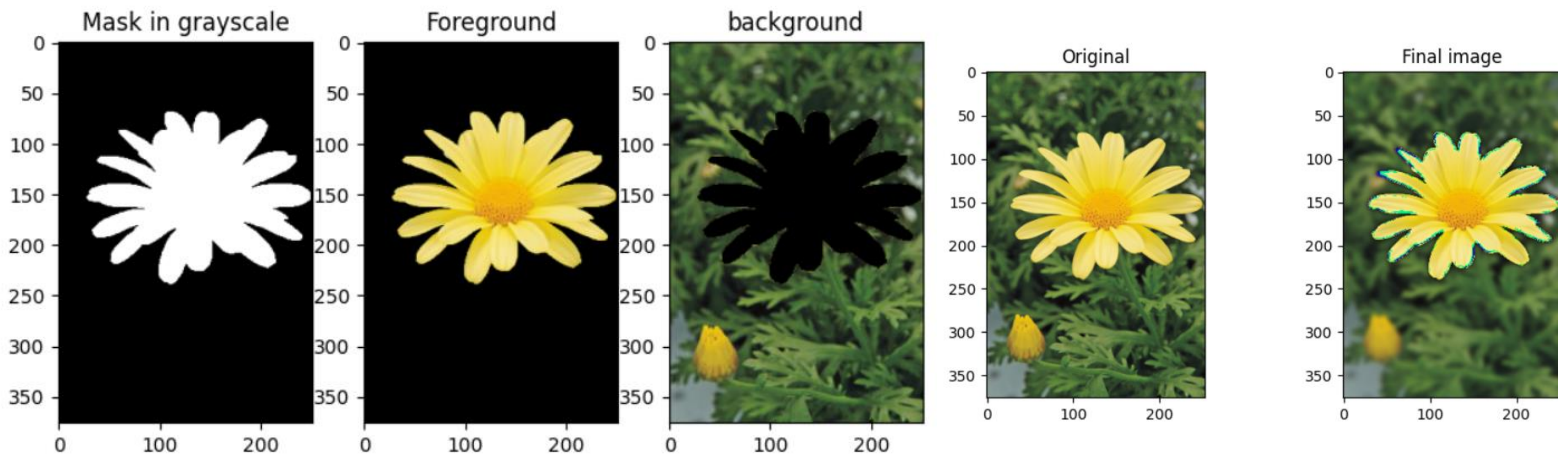
```
rect = (20,50,250,200)
bgdModel = np.zeros((1,65),np.float64)
fgdModel = np.zeros((1,65),np.float64)

mask = np.zeros(img.shape[:2],np.uint8)

cv.grabCut(img,mask,rect,bgdModel,fgdModel,5,cv.GC_INIT_WITH_RECT)

mask2 = np.where((mask==2)|(mask==0),0,1).astype('uint8')
# plot mask
plt.figure(figsize=(8,4))
plt.subplot(1, 3, 1)
plt.imshow(mask2, cmap='gray')
plt.title('Mask in grayscale')

foreground = img * mask2[:, :, np.newaxis]
background = img - foreground
blurred_background = cv.GaussianBlur(background, (15, 15),0)
final_image = blurred_background + foreground
```



Q9. (c) The dark (black) pixels are added to the blurred image when the foreground image and the blurred background image are layered on top of one another. Consequently, the flower's margin is a little bit black.